



Network Flow: Running Time

Max-Flow: Edmonds-Karp Algorithm, Dinitz's Algorithm
Max Bipartite Matching: Hopcroft-Karp-Karzanov algorithm

Residual Network G^f

- Given $G = (V, E)$, c , and a flow f
- $G^f = (V, E^f)$ with capacity c^f
- $(u, v) \in E^f$ if one of the followings holds
 - $(u, v) \in E$ and $f(u, v) < c(u, v)$: $c^f(u, v) = c(u, v) - f(u, v)$
 - $(v, u) \in E$ and $f(v, u) > 0$: $c^f(u, v) = f(v, u)$



Last Lecture – Ford-Fulkerson Method

- Always terminates for integer/rational capacities
- Not guaranteed to terminate for irrational capacities
- Time complexity for integer capacities: $O(|E| \cdot v(f_{\max}))$
 - not a polynomial time


Does the algorithm always halt?

- How about possibly irrational capacities?
- No, the algorithm do not always halt!

Non-terminating example [\[edit\]](#)

Consider the flow network shown on the right, with source s , sink t , capacities of edges e_1 , e_2 and e_3 respectively 1, $r = (\sqrt{5} - 1)/2$ and 1 and the capacity of all other edges some integer $M \geq 2$. The constant r was chosen so, that $r^2 = 1 - r$. We use augmenting paths according to the following table, where $p_1 = \{s, v_4, v_3, v_2, v_1, t\}$, $p_2 = \{s, v_2, v_3, v_4, t\}$ and $p_3 = \{s, v_1, v_2, v_3, t\}$.

Step	Augmenting path	Sent flow	Residual capacities		
			e_1	e_2	e_3
0			$r^0 = 1$	r	1
1	$\{s, v_2, v_3, t\}$	1	r^0	r^1	0
2	p_1	r^1	r^2	0	r^1
3	p_2	r^1	r^2	r^1	0
4	p_1	r^2	0	r^3	r^2
5	p_3	r^2	r^2	r^3	0



Note that after step 1 as well as after step 5, the residual capacities of edges e_1 , e_2 and e_3 are in the form r^n , r^{n+1} and 0, respectively, for some $n \in \mathbb{N}$. This means that we can use augmenting paths p_1 , p_2 , p_1 and p_3 infinitely many times and residual capacities of these edges will always be in the same form. Total flow in the network after step 5 is $1 + 2(r^1 + r^2)$. If we continue to use augmenting paths as above, the total flow converges to $1 + 2 \sum_{i=1}^{\infty} r^i = 3 + 2r$. However, note that there is a flow of value $2M + 1$, by sending M units of flow along sv_1t , 1 unit of flow along sv_2v_3t , and M units of flow along sv_4t . Therefore, the algorithm never terminates and the flow does not even converge to the maximum flow.^[4]


Another non-terminating example based on the [Euclidean algorithm](#) is given by [Backman & Huynh \(2018\)](#), where they also show that the worst case running-time of the Ford-Fulkerson algorithm on a network $G(V, E)$ in [ordinal numbers](#) is $\omega^{\Theta(|E|)}$.

Time Complexity?

- Assume all capacities are integers, what is the time complexity?
- Each iteration requires $O(|E|)$ time:
 - $O(|E|)$ is sufficient for finding p , updating f and G^f
- There are at most f_{max} iterations.
- Overall: $O(|E| \cdot f_{max})$
- Can we analyze it better?

Time Complexity?

- Can we analyze it better?
- It depends on how you choose p in each iteration!
- The complexity bound $O(|E| \cdot f_{max})$ is tight if choices of p are not carefully specified!



Method vs Algorithm

- Different choices of augmenting paths p give different implementation of Ford-Fulkerson.
- The description of Ford-Fulkerson Algorithm is incomplete.
- For this reason, it is sometimes called Ford-Fulkerson **Method**.
- Next Lecture Preview: Edmonds-Karp Algorithm, which implement Ford-Fulkerson Method with time complexity $O(|V| \cdot |E|^2)$.

Edmonds-Karp Algorithm


Edmonds-Karp Algorithm

EdmondsKarp($G = (V, E), s, t, c$):

1. initialize f such that $\forall e \in E: f(e) = 0$; initialize $G^f \leftarrow G$;
2. **while** there is an s - t path on G^f :
3. **find such a path p by BFS;**
4. find an edge $e \in p$ with minimum capacity b ;
5. update f that pushes b units of flow along p ;
6. update G^f ;
7. **endwhile**
8. **return** f

Why BFS?


- BFS maintains the distances
 - distance: num of edges, not weighted distance



A path found by an iteration of Edmonds-Karp Algorithm


Examples

- Can we choose the green path?




Examples

- We choose the red path!



Why BFS?


- In the residual network G^f , a **new appeared edge** can only go from a vertex at distance $t + 1$ to a vertex at distance t .
- Addition of such edges does not **decrease** the distance between s and u for every $u \in V$.
- **[Key Observation]** $\text{dist}(u)$ in G^f is non-decreasing.



The updates to the edges in G^f

Examples


- We choose the red path!



Weak Monotonicity to Strong Monotonicity

- $\text{dist}(u)$ can only be one of $0, 1, 2, \dots, |V|, \infty$
 - It can only be increased for $|V| + 1$ times!
- It's great that BFS buys us distance monotonicity!
- However, **weak** monotonicity is not enough.
- To make a progress, we need $\text{dist}(u)$ **strictly** increases for some $u \in V$, so that we can upper bound the number of iterations.

Counterexample: $\text{dist}(u)$ for all vertices remain unchanged after an iteration.




Towards Strong Monotonicity...

- Observation: At least one edge (u, v) on p is **saturated**, and this edge will be deleted in the next iteration.
- Each iteration will remove an edge from a vertex at distance i to a vertex at distance $i + 1$.
- Intuitively, we **cannot keep removing** such edges while keeping the distances of all vertices unchanged.

Towards Strong Monotonicity

- Suppose we are at the $(i + 1)$ -th iteration. f_i is the current flow, and p is the path found in G^{f_i} at the $(i + 1)$ -th iteration.
- We say that an edge (u, v) is **critical** if the amount of flow pushed along p is $c^{f_i}(u, v)$.
- A **critical** edge disappears in $G^{f_{i+1}}$.




Towards Strong Monotonicity

- Suppose we are at the $(i + 1)$ -th iteration. f_i is the current flow, and p is the path found in G^{f_i} at the $(i + 1)$ -th iteration.
- We say that an edge (u, v) is **critical** if the amount of flow pushed along p is $c^{f_i}(u, v)$.
- A **critical** edge disappears in $G^{f_{i+1}}$, but it may reappear in the future...
- We will try to bound the number of times (u, v) becomes critical.

Between two "critical"

A flow along p in G^{f_i} where (u, v) becomes critical




In $G^{f_{i+1}}$, (u, v) disappears, and (v, u) appears.




Before the next time (u, v) becomes critical again, (u, v) must first reappear!

Before (u, v) reappears, the algorithm must have found p going through (v, u) .



Between two "critical"



- Distance monotonicity: $\text{dist}^{i+j}(v) \geq \text{dist}^i(v)$.
- Thus, $\text{dist}^{i+j}(u) = \text{dist}^{i+j}(v) + 1 \geq \text{dist}^i(v) + 1 \geq \text{dist}^i(u) + 2$.
- The distance of u from s increases by 2 between two critical.

Putting Together


- The distance of u from s increases by 2 between two "critical".
- Distance takes value from $\{0, 1, \dots, |V|, \infty\}$, and never decrease.
- Thus, each edge can only be critical for $O(|V|)$ times.
- At least 1 edge become critical in one iteration.
- Total number of iterations is $O(|V| \cdot |E|)$.
- Each iteration takes $O(|E|)$ time.
- Overall time complexity for Edmonds-Karp: $O(|V| \cdot |E|^2)$.
- It can handle the issue with irrational numbers!

Can we improve?

- Learn from proof.
- $dist[u]$ is non-decreasing but not **strictly increasing**.
- Can we try to make some $dist$ strictly increasing?
- What if $dist[t]$ is strictly increasing?
 - $O(|V|)$ rounds increasing.

How to make it?

- We want to increase $dist[t]$.
- We should remove all shortest $s \rightarrow t$ path but not only one.





Dinic's Algorithm (Dinitz's Algorithm)

- Proposed by Yefim Dinitz (Soviet→ Israeli), in 1970.
 - Independent on Edmonds–Karp (1972).
- Updated by Shimon Even (Israeli). and Alon Itai (Israeli).
- Even gave lectures on "Dinic's algorithm", mispronouncing the name of the author while popularizing it.
- Time complexity: $O(|V|^2 \cdot |E|)$
 - Edmonds-Karp: $O(|V| \cdot |E|^2)$


Dinic's Algorithm – high-level ideas

- Build a **level graph**:
 - Vertices at Level i are at distance i .
 - Only edges go from a level to the next level are kept.
 - Can be done in $O(|E|)$ time using a similar idea to BFS.




Dinic's Algorithm – high-level ideas

- Find a **blocking flow** on the **level graph**:
 - Push flow on multiple s - t paths.
 - Each s - t path must contain a critical edge!




Dinic's Algorithm – high-level ideas

- Find a **blocking flow** on the **level graph**:
 - Push flow on multiple s - t paths.
 - Each s - t path must contain a critical edge!



a blocking flow



not a blocking flow: path $s \rightarrow a \rightarrow b \rightarrow t$ contains no critical edge

Dinic's Algorithm – Overview

- Initialize f to be the empty flow and $G^f = G$.
- Iteratively do the followings until $\text{dist}(t) = \infty$:
 - Construct the level graph G_L^f for G^f .
 - Find a blocking flow on G_L^f .
 - Update f and G^f .


Questions Remain

1. How many iterations do we need before termination?
2. How do we find a blocking flow?


How to find a block flow?

Dinic's Algorithm – high-level ideas

- Find a **blocking flow** on the **level graph**:
 - Push flow on multiple s - t paths.
 - Each s - t path must contain a critical edge!



a blocking flow




not a blocking flow: path s-a-b-t
contains no critical edge

Finding a blocking flow in a level graph...


Iteratively do the followings, until no path from s to t :

- Run a frontward DFS.
- Two possibilities:
 - End up at t : in this case, we update f (by pushing flow along the path) and remove the critical edge
 - End up at a dead-end, a vertex v with no out-going edges in G_L^f : in this case, we remove all the incoming edges of v


Sample Run.




Sample Run.




Sample Run.




Sample Run.




Sample Run.




Sample Run.



Sample Run.



Sample Run.



Finding a blocking flow in a level graph...


- At least one edge is removed after each search.
 - Total number of searches: $O(|E|)$
- Each search takes at most $|V|$ steps.
- Time complexity for each iteration of Dinic's algorithm: $O(|V| \cdot |E|)$.

Questions Remain

1. How do we find a blocking flow? ✓
2. How many iterations do we need before termination?


Simple yet important observations

- In the level graph $G_L^{f_i}$ at every iteration i , every s - t path has length $\text{dist}^i(t)$.
- Every shortest s - t path in G^{f_i} also appears in $G_L^{f_i}$.



Distance Monotonicity

- After one iteration, a new edge (u, v) appearing in $G^{f_{i+1}}$ (but not in G^{f_i}) must be "backward": $\text{dist}^i(u) = \text{dist}^i(v) + 1$.



Distance Monotonicity


- After one iteration, a new edge (u, v) appearing in $G^{f_{i+1}}$ (but not in G^{f_i}) must satisfy $\text{dist}^i(u) = \text{dist}^i(v) + 1$.
- Such additions of edges cannot reduce the distance for any vertex!
- We again have that $\text{dist}(u)$ is non-decreasing!
- Can we have strong monotonicity?

Taking a closer look...

- All the paths in $G_L^{f^i}$ with length $\text{dist}^i(t)$ are “blocked” after the i -th iteration.
- Thus, a path in the $(i + 1)$ -th iteration must use some edges **that are not in $G_L^{f^i}$** .


Taking a closer look...

- This new edge may be a "backward" edge whose reverse was a critical edge in the previous iteration.
- In this case, $\text{dist}(t)$ is increased by at least 2.



Taking a closer look...

- Or it may be an edge in G^{f_i} , but not in $G_L^{f_i}$.
- In this case, $\text{dist}(t)$ is increased by at least 1.



Taking a closer look...

- In both cases: $\text{dist}^{i+1}(t) > \text{dist}^i(t)$
- Let's prove it rigorously then...

Proving $\text{dist}^{i+1}(t) > \text{dist}^i(t)$

- Consider an arbitrary s - t path p in $G_L^{f^{i+1}}$ with length $\text{dist}^{i+1}(t)$.
- We have $\text{dist}^{i+1}(t) \geq \text{dist}^i(t)$ by monotonicity.
- Suppose for the sake of contradiction that $\text{dist}^{i+1}(t) = \text{dist}^i(t)$.
- Case 1: all edges in p also appear in $G_L^{f^i}$
- Then p is a shortest path containing no critical edges in $G_L^{f^i}$
- Contradicting to the definition of blocking flow!

Proving $\text{dist}^{i+1}(t) > \text{dist}^i(t)$

- Case 2: p contains an edge (u, v) that is not in G_L^{fi}
- If (u, v) was not in G^{fi} , then (v, u) was critical in the last iteration. We have $\text{dist}^i(u) = \text{dist}^i(v) + 1$.
- If (u, v) was in G^{fi} but not G_L^{fi} , by the definition of level graph, we have $\text{dist}^i(u) \geq \text{dist}^i(v)$.
- In both cases above, $\text{dist}^i(u) \geq \text{dist}^i(v)$.
- We have $\text{dist}^{i+1}(u) \geq \text{dist}^i(u)$ by monotonicity,
- and we have $\text{dist}^{i+1}(v, t) \geq \text{dist}^i(v, t)$. (why?)

Proving $\text{dist}^{i+1}(t) > \text{dist}^i(t)$

- Case 2: p contains an edge (u, v) that is not in $G_L^{f_i}$
- Fact i: $\text{dist}^i(u) \geq \text{dist}^i(v)$.
- Fact ii: $\text{dist}^{i+1}(u) \geq \text{dist}^i(u)$.
- Fact iii: $\text{dist}^{i+1}(v, t) \geq \text{dist}^i(v, t)$.

Putting together:

$$\begin{aligned}\text{dist}^{i+1}(t) &= \text{dist}^{i+1}(u) + 1 + \text{dist}^{i+1}(v, t) \\ &\geq \text{dist}^i(u) + 1 + \text{dist}^i(v, t) && \text{(Fact ii and iii)} \\ &\geq \text{dist}^i(v) + 1 + \text{dist}^i(v, t) && \text{(Fact i)} \\ &\geq \text{dist}^i(t) + 1 && \text{(triangle inequality)}\end{aligned}$$



Putting Together...

- $\text{dist}(t)$ is increased by at least 1 after each iteration.
- $\text{dist}(t)$ takes value from $\{0, 1, \dots, |V|, \infty\}$, so it can be increased for at most $O(|V|)$ times.
- Total number of iterations is at most $O(|V|)$.

Other Algorithms for Max-Flow

- Improvements to Dinic's algorithm:
 - [Malhotra, Kumar & Maheshwari, 1978]: $O(|V|^3)$
 - Dynamic tree: $O(|V| \cdot |E| \cdot \log|V|)$
- Push-relabel algorithm [Goldberg & Tarjan, 1988]
 - $O(|V|^2|E|)$, later improved to $O(|V|^3)$, $O(|V|^2\sqrt{|E|})$, $O(|V||E|\log\frac{|V|^2}{|E|})$
- [King, Rao & Tarjan, 1994] and [Orlin, 2013]: $O(|V| \cdot |E|)$
- Interior-point-method-based algorithms:
 - [Kathuria, Liu & Sidford, 2020] $|E|^{\frac{4}{3}+o(1)}U^{\frac{1}{3}}$
 - [BLNPSSSW, 2020] [BLLSSSW, 2021] $\tilde{O}\left((|E| + |V|^{\frac{3}{2}})\log U\right)$
 - [Gao, Liu & Peng, 2021] $\tilde{O}\left(|E|^{\frac{3}{2}-\frac{1}{328}}\log U\right)$

Questions Remain

1. How do we find a blocking flow? 
2. How many iterations do we need before termination? 

Overall Time Complexity for Dinic's Algorithm


- Each iteration: $O(|V| \cdot |E|)$.
- We need at most $O(|V|)$ iterations.
- Overall time complexity for Dinic's algorithm: $O(|V|^2 \cdot |E|)$.

Hopcroft–Karp–Karzanov algorithm

- Find a maximum bipartite matching in $O(|E| \cdot \sqrt{|V|})$ time.
- Proposed independently by Hopcroft-Karp and Karzanov.
- Can be viewed as a special case of Dinic's algorithm.


Conversion to Max-Flow Problem

Set the capacity to 1 for all edges.



Conversion to Max-Flow Problem

Dinic's algorithm runs in $O\left(|E| \cdot \sqrt{|V|}\right)$ time for this special case.



Conversion to Max-Flow Problem


- Integrality theorem also holds for Dinic's algorithm:
 - The flow output by Dinic's algorithm in our case is integral.
- We aim to show Dinic's algorithm runs in $O(|E| \cdot \sqrt{|V|})$ time.
- Step 1: Finding a blocking flow in a level graph takes $O(|E|)$ time.
- Step 2: Number of iterations is at most $2\sqrt{|V|}$.

Step 1: Finding a blocking flow in a level graph takes $O(|E|)$ time.


Iteratively do the followings, until no path from s to t :

- Perform **DFS** from s
- If we reach t , delete **all** edges on the s - t path (why can we do this?) **and start over from s** .
- If we ever go backward, delete the edge just travelled. (why can we do this?)


Step 1: Finding a blocking flow in a level graph takes $O(|E|)$ time.




Step 1: Finding a blocking flow in a level graph takes $O(|E|)$ time.



Step 1: Finding a blocking flow in a level graph takes $O(|E|)$ time.




Step 1: Finding a blocking flow in a level graph takes $O(|E|)$ time.




Step 1: Finding a blocking flow in a level graph takes $O(|E|)$ time.

An s - t path is found, remove all edges from the path.




Step 1: Finding a blocking flow in a level graph takes $O(|E|)$ time.

Start over...




Step 1: Finding a blocking flow in a level graph takes $O(|E|)$ time.




Step 1: Finding a blocking flow in a level graph takes $O(|E|)$ time.

We have to go backward now; delete the edge just travelled.




Step 1: Finding a blocking flow in a level graph takes $O(|E|)$ time.




Step 1: Finding a blocking flow in a level graph takes $O(|E|)$ time.

Again, we have to go backward; delete the edge just travelled.




Step 1: Finding a blocking flow in a level graph takes $O(|E|)$ time.

Again, we have to go backward; delete the edge just travelled.




Step 1: Finding a blocking flow in a level graph takes $O(|E|)$ time.


Again, we have to go backward; delete the edge just travelled.



Step 1: Finding a blocking flow in a level graph takes $O(|E|)$ time.




Step 1: Finding a blocking flow in a level graph takes $O(|E|)$ time.




Step 1: Finding a blocking flow in a level graph takes $O(|E|)$ time.

Find another s - t path; delete all edges on the path




Step 1: Finding a blocking flow in a level graph takes $O(|E|)$ time.

We are done!



Step 1: Finding a blocking flow in a level graph takes $O(|E|)$ time.

We have obtained a blocking flow!



Step 1: Finding a blocking flow in a level graph takes $O(|E|)$ time.

Time complexity: $O(|E|)$


- Each edge is visited at most once.

Step 2: Number of iterations is at most $2\sqrt{|V|}$.

- If the algorithm terminates within $\sqrt{|V|}$ iterations, we are already done!
- Otherwise, let f be the flow after $\sqrt{|V|}$ iterations.
- Claim: the maximum flow in G^f has value at most $\sqrt{|V|}$.
- If the claim is true, we can stop after another $\sqrt{|V|}$ rounds.

Observation on G^f

- In each iteration, for each $v \in V \setminus \{s, t\}$, either its in-degree is 1, or its out-degree is 1.
- Proof. By Induction...
- At the beginning, this is clearly true.




Observation on G^f

- In each iteration, for each $v \in V \setminus \{s, t\}$, either its in-degree is 1, or its out-degree is 1.
- Proof. At the beginning, this is clearly true.
- For each iteration, the amount of flow going through v is either 0 or 1.
- If it is 0, v 's in-degree and out-degree are unchanged.
- Otherwise, exactly one in-edge and one out-edge are flipped; the property is still maintained.





Let us check!



The maximum flow in G^f has value at most $\sqrt{|V|}$

- Integrality Theorem: there exists a maximum integral flow f' in G^f .
- f' consists of **vertex-disjoint** paths with flow 1!



The maximum flow in G^f has value at most $\sqrt{|V|}$

- Max-flow on G^f, f' , is integral and consists of edge-disjoint paths.
- By our analysis to Dinic's algorithm, $\text{dist}^{G^f}(s, t) \geq \sqrt{|V|}$.
- Each path in f' has length at least $\sqrt{|V|}$.
- There are at most $\frac{|V|}{\sqrt{|V|}} = \sqrt{|V|}$ paths in f' by vertex-disjointness.
- $v(f') \leq \sqrt{|V|}$

Step 2: Number of iterations is at most $2\sqrt{|V|}$.

- If the algorithm terminates within $\sqrt{|V|}$ iterations, we are already done!
- Otherwise, let f be the flow after $\sqrt{|V|}$ iterations.
- Claim: the maximum flow in G^f has value at most $\sqrt{|V|}$. ✓
- Each iteration increase the value of flow by at least 1.
- Thus, the algorithm will terminate within at most another $\sqrt{|V|}$ iterations.
- Total number of iterations: $2\sqrt{|V|}$.

Putting Together...

- Step 1: Finding a blocking flow in a level graph takes $O(|E|)$ time.
- Step 2: Number of iterations is at most $2\sqrt{|V|}$.
- Overall time complexity: $O(|E| \cdot \sqrt{|V|})$



Today's Lecture

Maximum Flow Problem:

- Edmonds-Karp Algorithm
 - Implement Ford-Fulkerson method by BFS
 - $O(|V| \cdot |E|^2)$
- Dinic's Algorithm
 - Push flow on multiple paths at one iteration
 - Level graph and blocking flow
 - $O(|V|^2 \cdot |E|)$

Maximum Bipartite Matching Problem:

- Hopcroft-Karp-Karzanov algorithm
 - Apply Dinic's algorithm
 - $O(|E| \cdot \sqrt{|V|})$